

Peer-graded Assignment: Prediction Assignment Writeup

2023-10-07

Summary

This document is the final report of the Peer Assessment project from the Practical Machine Learning course, which is a part of the Coursera John's Hopkins University Data Science Specialization. It was written and coded in RStudio, using its knitr functions and published in the html and markdown format. The goal of this project is to predict the manner in which the six participants performed the exercises. The machine learning algorithm, which uses the classe variable in the training set, is applied to the 20 test cases available in the test data.

Introduction

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement – a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, your goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways.

Data Source

The training and test data for this project are collected using the link below:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>

The data for this project come from this source: <http://groupware.les.inf.puc-rio.br/har>.

The full reference of this data is as follows:

Velloso, E.; Bulling, A.; Gellersen, H.; Ugulino, W.; Fuks, H. "Qualitative Activity Recognition of Weight Lifting Exercises. Proceedings of 4th International Conference in Cooperation with SIGCHI (Augmented Human '13)". Stuttgart, Germany: ACM SIGCHI, 2013.

Loading and Cleaning of Data

Set working directory.

```
setwd("~/Downloads")
```

Load required R packages and set a seed.

```
library(lattice)
library(ggplot2)
library(caret)
library(rpart)
library(rpart.plot)
library(corrplot)

## corrplot 0.92 loaded

library(rattle)

## Loading required package: tibble

## Loading required package: bitops

## Rattle: A free graphical interface for data science with R.
## Version 5.5.1 Copyright (c) 2006-2021 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.

library(randomForest)

## randomForest 4.7-1.1

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:rattle':
##
##     importance

## The following object is masked from 'package:ggplot2':
##
##     margin

library(RColorBrewer)

set.seed(222)
```

Load data for training and test datasets.

```
url_train <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-
training.csv"
url_quiz <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-
testing.csv"

data_train <- read.csv(url(url_train), strip.white = TRUE, na.strings =
c("NA", ""))
data_quiz <- read.csv(url(url_quiz), strip.white = TRUE, na.strings =
c("NA", ""))
```

```
dim(data_train)
## [1] 19622  160
dim(data_quiz)
## [1]  20 160
```

Crete two partitions (75% and 25%) within the original training dataset.

```
in_train  <- createDataPartition(data_train$classe, p=0.75, list=FALSE)
train_set <- data_train[ in_train, ]
test_set  <- data_train[-in_train, ]

dim(train_set)
## [1] 14718  160
dim(test_set)
## [1] 4904  160
```

The two datasets (train_set and test_set) have a large number of NA values as well as near-zero-variance (NZV) variables. Both will be removed together with their ID variables.

```
nzv_var <- nearZeroVar(train_set)

train_set <- train_set[ , -nzv_var]
test_set  <- test_set [ , -nzv_var]

dim(train_set)
## [1] 14718  120
dim(test_set)
## [1] 4904  120
```

Remove variables that are mostly NA. A threshold of 95 % is selected.

```
na_var <- sapply(train_set, function(x) mean(is.na(x))) > 0.95
train_set <- train_set[ , na_var == FALSE]
test_set  <- test_set [ , na_var == FALSE]

dim(train_set)
## [1] 14718  59
dim(test_set)
## [1] 4904  59
```

Since columns 1 to 5 are identification variables only, they will be removed as well.

```
train_set <- train_set[ , -(1:5)]
test_set  <- test_set [ , -(1:5)]

dim(train_set)
## [1] 14718    54

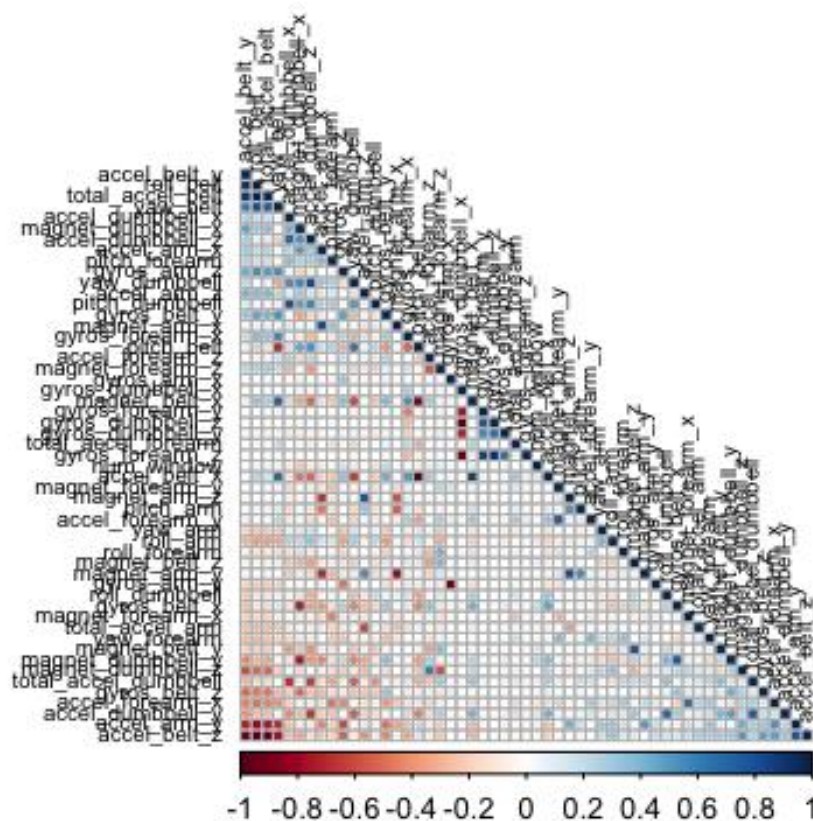
dim(test_set)
## [1] 4904     54
```

The number of variables for the analysis has been reduced from the original 160 down to 54.

Correlation Analysis

Correlation analysis between the variables before the modeling work itself is done. The “FPC” is used as the first principal component order.

```
corr_matrix <- cor(train_set[ , -54])
corrplot(corr_matrix, order = "FPC", method = "circle", type = "lower",
         tl.cex = 0.6, tl.col = rgb(0, 0, 0))
```



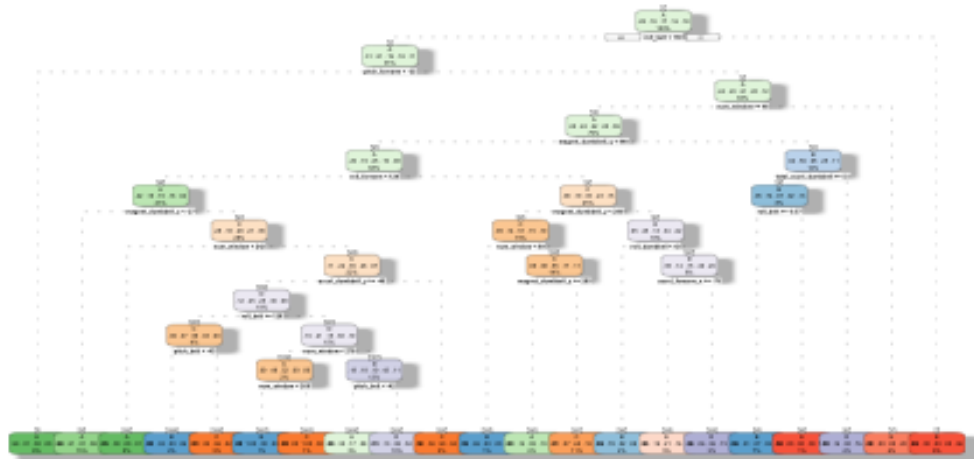
If two variables are highly correlated their colors are either dark blue (for a positive correlation) or dark red (for a negative correlations). Because there are only few strong correlations among the input variables, the Principal Components Analysis (PCA) will not be performed in this analysis. Instead, a few different prediction models will be built to have a better accuracy.

Prediction Models

Decision Tree Model

```
set.seed(2222)
fit_decision_tree <- rpart(classe ~ ., data = train_set, method="class")
fancyRpartPlot(fit_decision_tree)

## Warning: labs do not fit even at cex 0.15, there may be some overplotting
```



Rattle 2023-Oct-10 10:07:57 uma

Predictions of the decision tree model on test_set.

```
predict_decision_tree <- predict(fit_decision_tree, newdata = test_set,
type="class")
conf_matrix_decision_tree <- confusionMatrix(predict_decision_tree,
factor(test_set$classe))
conf_matrix_decision_tree
```

Confusion Matrix and Statistics

##

Reference

Prediction	A	B	C	D	E
A	1238	218	37	76	36
B	41	547	28	30	19
C	8	53	688	114	38
D	70	91	50	518	111
E	38	40	52	66	697

##

Overall Statistics

##

Accuracy : 0.752

95% CI : (0.7397, 0.7641)

No Information Rate : 0.2845

P-Value [Acc > NIR] : < 2.2e-16

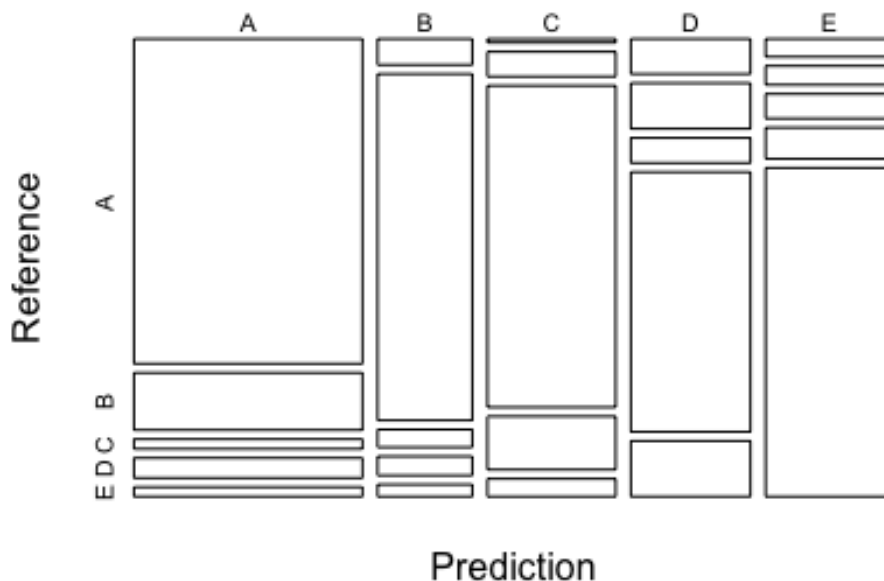
```
##
##          Kappa : 0.685
##
## McNemar's Test P-Value : < 2.2e-16
##
## Statistics by Class:
##
##          Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.8875  0.5764  0.8047  0.6443  0.7736
## Specificity      0.8954  0.9702  0.9474  0.9215  0.9510
## Pos Pred Value   0.7713  0.8226  0.7636  0.6167  0.7805
## Neg Pred Value    0.9524  0.9052  0.9583  0.9296  0.9491
## Prevalence       0.2845  0.1935  0.1743  0.1639  0.1837
## Detection Rate   0.2524  0.1115  0.1403  0.1056  0.1421
## Detection Prevalence 0.3273  0.1356  0.1837  0.1713  0.1821
## Balanced Accuracy 0.8914  0.7733  0.8760  0.7829  0.8623
```

The predictive accuracy of the decision tree model is relatively low at 75.2 %.

Plot the predictive accuracy of the decision tree model.

```
plot(conf_matrix_decision_tree$table, col =
conf_matrix_decision_tree$byClass,
     main = paste("Decision Tree Model: Predictive Accuracy =",
round(conf_matrix_decision_tree$overall['Accuracy'], 4)))
```

Decision Tree Model: Predictive Accuracy = 0.752



##

Generalised Boosted Model (GBM)

```
set.seed(2222)
ctrl_GBM <- trainControl(method = "repeatedcv", number = 5, repeats = 2)
fit_GBM <- train(classe ~ ., data = train_set, method = "gbm", trControl =
ctrl_GBM, verbose = FALSE)
fit_GBM$finalModel

## A gradient boosted model with multinomial loss function.
## 150 iterations were performed.
## There were 53 predictors of which 53 had non-zero influence.
```

Predictions of the GBM on test_set.

```
predict_GBM <- predict(fit_GBM, newdata = test_set)
conf_matrix_GBM <- confusionMatrix(predict_GBM, factor(test_set$classe))
conf_matrix_GBM

## Confusion Matrix and Statistics
##
##           Reference
## Prediction    A    B    C    D    E
##           A 1391    6    0    1    0
##           B   4  927    3    4    4
##           C   0   12  843   10    2
```



```
##           D      0      4      9  786    11
##           E      0      0      0    3  884
##
## Overall Statistics
##
##           Accuracy : 0.9851
##           95% CI : (0.9813, 0.9883)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9812
##
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9971  0.9768  0.9860  0.9776  0.9811
## Specificity      0.9980  0.9962  0.9941  0.9941  0.9993
## Pos Pred Value   0.9950  0.9841  0.9723  0.9704  0.9966
## Neg Pred Value   0.9989  0.9944  0.9970  0.9956  0.9958
## Prevalence       0.2845  0.1935  0.1743  0.1639  0.1837
## Detection Rate   0.2836  0.1890  0.1719  0.1603  0.1803
## Detection Prevalence 0.2851  0.1921  0.1768  0.1652  0.1809
## Balanced Accuracy 0.9976  0.9865  0.9900  0.9859  0.9902
```

The predictive accuracy of the GBM is relatively high at 98.57 %. ## Random Forest Model

```
set.seed(2222)
ctrl_RF <- trainControl(method = "repeatedcv", number = 5, repeats = 2)
fit_RF <- train(classe ~ ., data = train_set, method = "rf",
  trControl = ctrl_RF, verbose = FALSE)
fit_RF$finalModel

##
## Call:
## randomForest(x = x, y = y, mtry = param$mtry, verbose = FALSE)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 27
##
##           OOB estimate of error rate: 0.24%
## Confusion matrix:
##           A      B      C      D      E class.error
## A 4183      1      0      0      1 0.0004778973
## B   8 2836      3      1      0 0.0042134831
## C   0   6 2561      0      0 0.0023373588
## D   0   0   7 2404      1 0.0033167496
## E   0   1   0   7 2698 0.0029563932
```

Predictions of the random forest model on test_set.

```

predict_RF <- predict(fit_RF, newdata = test_set)
conf_matrix_RF <- confusionMatrix(predict_RF, factor(test_set$classe))
conf_matrix_RF

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    A    B    C    D    E
##      A 1395     3     0     0     0
##      B     0   946     2     0     0
##      C     0     0   853     6     0
##      D     0     0     0   798     1
##      E     0     0     0     0   900
##
## Overall Statistics
##
##              Accuracy : 0.9976
##              95% CI : (0.9957, 0.9987)
##      No Information Rate : 0.2845
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.9969
##
##      McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: A Class: B Class: C Class: D Class: E
## Sensitivity          1.0000  0.9968  0.9977  0.9925  0.9989
## Specificity          0.9991  0.9995  0.9985  0.9998  1.0000
## Pos Pred Value       0.9979  0.9979  0.9930  0.9987  1.0000
## Neg Pred Value       1.0000  0.9992  0.9995  0.9985  0.9998
## Prevalence           0.2845  0.1935  0.1743  0.1639  0.1837
## Detection Rate       0.2845  0.1929  0.1739  0.1627  0.1835
## Detection Prevalence 0.2851  0.1933  0.1752  0.1629  0.1835
## Balanced Accuracy    0.9996  0.9982  0.9981  0.9961  0.9994

```

The predictive accuracy of the Random Forest model is excellent at 99.8 %.

Applying the Best Predictive Model to the Test Data

The following are the predictive accuracy of the three models:

Decision Tree Model: 75.20 % Generalized Boosted Model: 98.57 % Random Forest Model: 99.80 % The Random Forest model is selected and applied to make predictions on the 20 data points from the original testing dataset (data_quiz).

```

predict_quiz <- as.data.frame(predict(fit_RF, newdata = data_quiz))
predict_quiz

```

```
## predict(fit_RF, newdata = data_quiz)
## 1 B
## 2 A
## 3 B
## 4 A
## 5 A
## 6 E
## 7 D
## 8 B
## 9 A
## 10 A
## 11 B
## 12 C
## 13 B
## 14 A
## 15 E
## 16 E
## 17 A
## 18 B
## 19 B
## 20 B
```